

```
In [25]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC, SVR
from sklearn.metrics import accuracy_score, mean_squared_error
import matplotlib.pyplot as plt
```

```
In [26]: # Load dataset
file_path = 'NFLX.csv'
data = pd.read_csv(file_path)
```

```
In [27]: # Display the first few rows of the dataset
print(data.head())
```

	Date	Open	High	Low	Close	Adj Close	\
0	2018-02-05	262.000000	267.899994	250.029999	254.259995	254.259995	
1	2018-02-06	247.699997	266.700012	245.000000	265.720001	265.720001	
2	2018-02-07	266.579987	272.450012	264.329987	264.559998	264.559998	
3	2018-02-08	267.079987	267.619995	250.000000	250.100006	250.100006	
4	2018-02-09	253.850006	255.800003	236.110001	249.470001	249.470001	

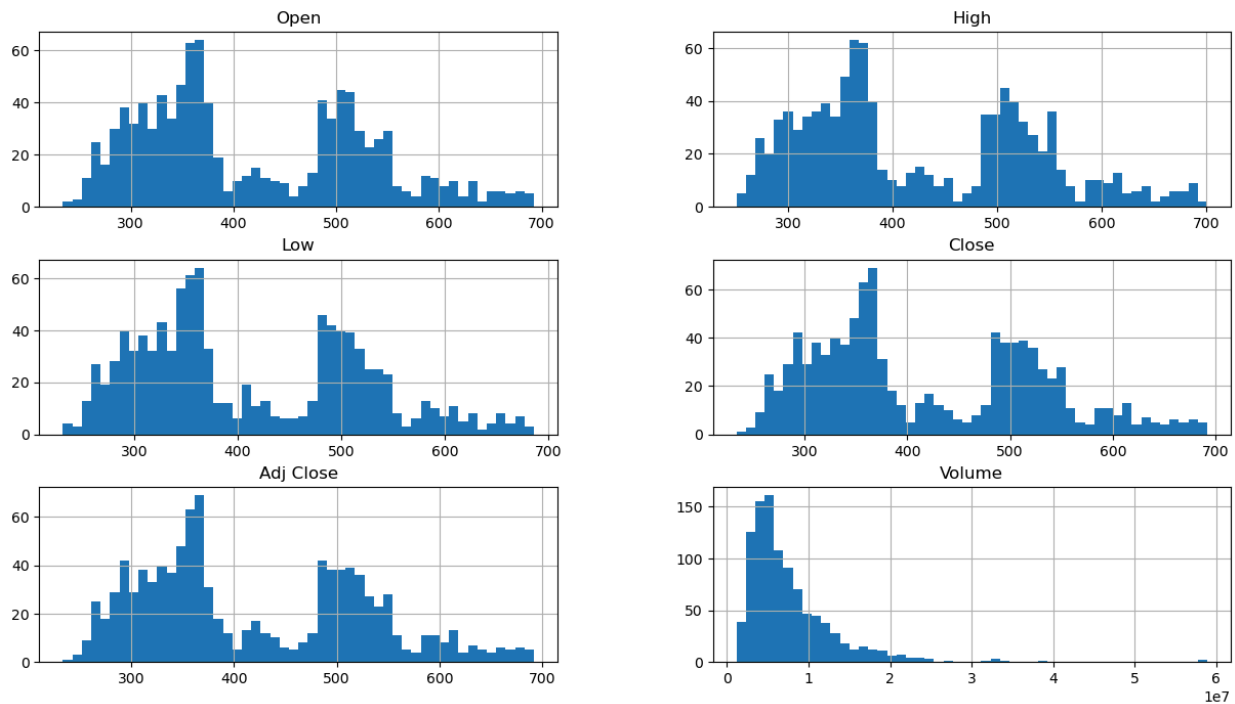
	Volume
0	11896100
1	12595800
2	8981500
3	9306700
4	16906900

```
In [28]: # Looking at the basic statistics about the data
display(data.describe())
```

	Open	High	Low	Close	Adj Close	Volume
count	1009.000000	1009.000000	1009.000000	1009.000000	1009.000000	1.009000e+03
mean	419.059673	425.320703	412.374044	419.000733	419.000733	7.570685e+06
std	108.537532	109.262960	107.555867	108.289999	108.289999	5.465535e+06
min	233.919998	250.649994	231.229996	233.880005	233.880005	1.144000e+06
25%	331.489990	336.299988	326.000000	331.619995	331.619995	4.091900e+06
50%	377.769989	383.010010	370.880005	378.670013	378.670013	5.934500e+06
75%	509.130005	515.630005	502.529999	509.079987	509.079987	9.322400e+06
max	692.349976	700.989990	686.090027	691.690002	691.690002	5.890430e+07

```
In [29]: # Looking at the data distributions

data.hist(bins=50, figsize=(15,8))
plt.show()
```



```
In [30]: # Check for duplicate observations
duplicate_count = data.duplicated().sum()
print(f"There are {duplicate_count} duplicates in the dataset")

# Delete duplicate data if found
if duplicate_count > 0:
    data.drop_duplicates(inplace=True)
    print("\nDuplicate observations were deleted.")
    # Print the shape of data after removal of duplicates
    print("Data shape after duplicate removal:{0}".format(data.shape))
```

There are 0 duplicates in the dataset

```
In [31]: # Gets number of unique values for each column
unique_values_per_attrib = data.nunique()

# Records columns to delete
single_value_columns = [i for i, value_count in enumerate(unique_values_per_attrib) if
    print("There are", len(single_value_columns), "single-valued columns in the dataset")

# Deletes single-value columns, if exist
if len(single_value_columns) > 0:
    data.drop(single_value_columns#, axis=1, inplace=True)
    print("\nSingle-valued columns were removed.")
    # Prints the shape of data after removal of single-value columns
    print("\nData shape after single-value column removal:", data.shape)
```

There are 0 single-valued columns in the dataset

```
In [32]: # Handle missing values
data = data.dropna()
```

```
In [33]: # Convert 'Date' to datetime format
data['Date'] = pd.to_datetime(data['Date'])
```

```
In [34]: # Sort by date
data = data.sort_values(by='Date')
```

```

In [35]: # Use 'Close' prices for prediction
X = data[['Close']].values
y = np.sign(np.diff(X.flatten())) # +1 for up, -1 for down, 0 for no change
y = np.concatenate([[0], y]) # Align y with X

In [36]: # Feature scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

In [37]: # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, random

In [38]: # Linear SVM Classification
linear_svm = SVC(kernel='linear')
linear_svm.fit(X_train, y_train)
y_pred_linear = linear_svm.predict(X_test)
print("Accuracy (Linear SVM):", accuracy_score(y_test, y_pred_linear))

Accuracy (Linear SVM): 0.44554455445544555

In [39]: # Soft Margin Classification
soft_margin_svm = SVC(kernel='linear', C=0.1)
soft_margin_svm.fit(X_train, y_train)
y_pred_soft = soft_margin_svm.predict(X_test)
print("Accuracy (Soft Margin SVM):", accuracy_score(y_test, y_pred_soft))

Accuracy (Soft Margin SVM): 0.44554455445544555

In [40]: # Nonlinear SVM Classification with Polynomial Kernel
poly_svm = SVC(kernel='poly', degree=3, C=1)
poly_svm.fit(X_train, y_train)
y_pred_poly = poly_svm.predict(X_test)
print("Accuracy (Polynomial Kernel SVM):", accuracy_score(y_test, y_pred_poly))

Accuracy (Polynomial Kernel SVM): 0.44554455445544555

In [41]: # Nonlinear SVM Classification with Gaussian RBF Kernel
rbf_svm = SVC(kernel='rbf', gamma='scale', C=1)
rbf_svm.fit(X_train, y_train)
y_pred_rbf = rbf_svm.predict(X_test)
print("Accuracy (RBF Kernel SVM):", accuracy_score(y_test, y_pred_rbf))

Accuracy (RBF Kernel SVM): 0.5148514851485149

In [42]: # Preparing data for SVR
X_prices = X_scaled[:-1] # Use all but the last day
y_prices = X_scaled[1:, 0] # Predict the closing price for the next day

In [43]: # Hyperparameter tuning for SVR
svr_param_grid = {
    'kernel': ['linear', 'poly', 'rbf'],
    'C': [0.1, 1, 10],
    'degree': [2, 3, 4],
    'gamma': ['scale', 'auto'],
    'epsilon': [0.01, 0.1, 0.2]
}
best_svr = GridSearchCV(SVR(), svr_param_grid, cv=5, n_jobs=-1, scoring='neg_mean_squa

```

```
best_svr.fit(X_prices, y_prices)
best_svr_model = best_svr.best_estimator_
```

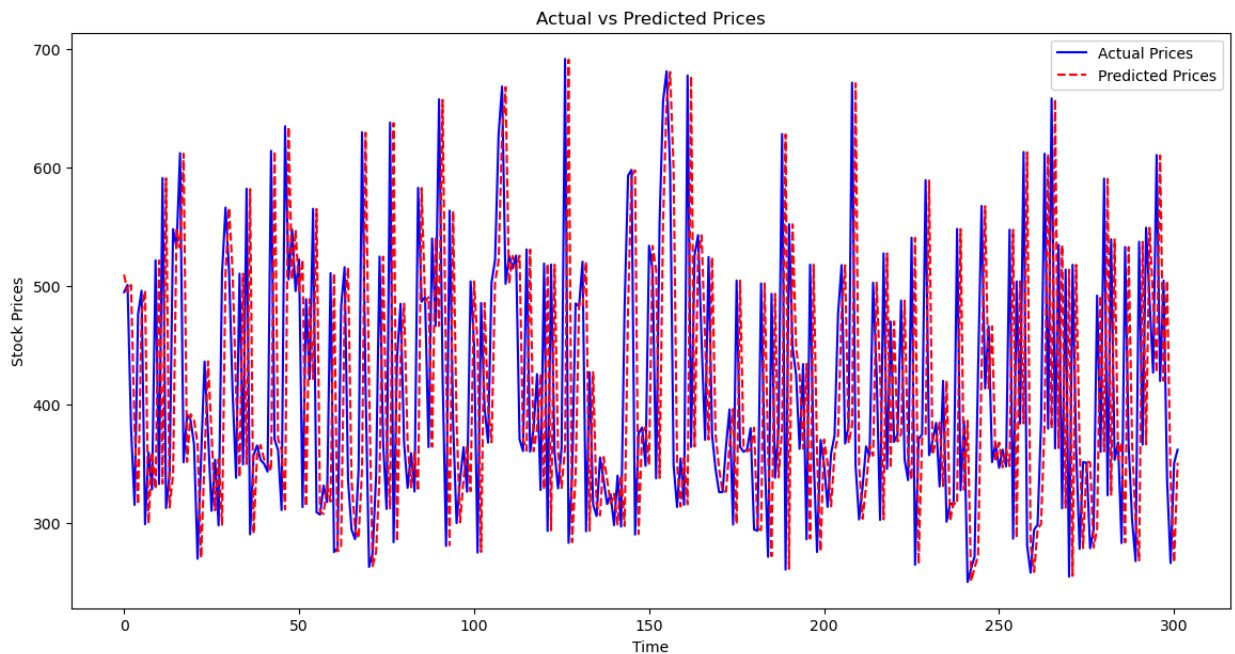
```
In [44]: # SVM Regression for Predicting Prices
best_svr_model.fit(X_prices, y_prices)
```

```
Out[44]: SVR
SVR(C=10, degree=2, epsilon=0.01, kernel='linear')
```

```
In [45]: # Predicting the test set
predicted_prices = best_svr_model.predict(X_test[:-1])
mse = mean_squared_error(X_test[1:, 0], predicted_prices)
print("MSE (Best SVR):", mse)
```

MSE (Best SVR): 1.9008962177167037

```
In [46]: # Visualizing Actual vs Predicted Prices
plt.figure(figsize=(14, 7))
plt.plot(scaler.inverse_transform(X_test[1:])[0], color='blue', label='Actual Prices')
plt.plot(scaler.inverse_transform(np.hstack([predicted_prices.reshape(-1, 1), np.zeros(
plt.title('Actual vs Predicted Prices')
plt.xlabel('Time')
plt.ylabel('Stock Prices')
plt.legend()
plt.show()
```



```
In [ ]:
```

```
In [ ]:
```